

EE559-HW7-code-oconnort-6038881588

July 22, 2024

Tim O'Connor

603-888-1588

oconnort@usc.edu

EE559 - Rajati - Summer 2024

HW 7 - Programming Part

4. Programming Assignment: Parkinsons Telemonitoring

- (a) Download the Parkinsons Telemonitoring Data Set from: <http://archive.ics.scu.edu/ml/datasets/Parkinsons+Telemonitoring>. Choose 70% of the data randomly as the training set.
- (b) Use metric learning with Gaussian kernels to estimate each of the outputs motor UPDRS and total UPDRS from the features. As metric learning uses a low dimensional transformation of the features except the non-predictive feature subject#, use 5-fold cross-validation to decide the number of components from $M=5,10,15,p$, where p is the number of all predictive features you can use. Initialize the linear transformation with PCA features for $M=5,10,15$ and with original features for $M=p$. This corresponds to setting `intit` as (default='auto'). Remember to standardize your features. Report the R^2 on training and test sets for each of the outputs. (30 pts)

```
[92]: import pandas as pd
from sklearn.model_selection import train_test_split, KFold
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.kernel_ridge import KernelRidge
from sklearn.metrics import r2_score
from metric_learn import MLKR
from joblib import Parallel, delayed

path = './parkinsons_updrs.data'

data = pd.read_csv(path)

X = data.drop(columns=['subject#', 'motor_UPDRS', 'total_UPDRS'])
y_motor = data['motor_UPDRS']
y_total = data['total_UPDRS']
```

```

X_train, X_test, y_motor_train, y_motor_test, y_total_train, y_total_test = ␣
    ↪ train_test_split(X, y_motor, y_total, test_size=0.3, random_state=42)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

kf = KFold(n_splits=5, shuffle=True, random_state=42)

n = 5 #originally tried to run all of the possible components in a loop, but it ␣
    ↪ was too computationally expensive, 5 was able to finish.

pca_transformations = PCA(n_components=n).fit(X_train_scaled)

def process_fold(train_i, value_i, n, X_train_scaled, y_motor_train, ␣
    ↪ y_total_train):
    X_train_kf, X_value_kf = X_train_scaled[train_i], X_train_scaled[value_i]
    y_motor_train_kf, y_motor_value_kf = y_motor_train.iloc[train_i], ␣
    ↪ y_motor_train.iloc[value_i]
    y_total_train_kf, y_total_value_kf = y_total_train.iloc[train_i], ␣
    ↪ y_total_train.iloc[value_i]

    X_motor_train_PCA = pca_transformations.transform(X_train_kf)
    X_motor_test_PCA = pca_transformations.transform(X_value_kf)

    mlkr_motor = MLKR(n_components=n, init='pca')
    mlkr_motor.fit(X_motor_train_PCA, y_motor_train_kf)
    X_motor_train_MLKR = mlkr_motor.transform(X_motor_train_PCA)
    X_motor_test_MLKR = mlkr_motor.transform(X_motor_test_PCA)

    gauss_kernel_motor = KernelRidge(kernel='rbf')
    gauss_kernel_motor.fit(X_motor_train_MLKR, y_motor_train_kf)

    y_train_pred_motor = gauss_kernel_motor.predict(X_motor_train_MLKR)
    y_test_pred_motor = gauss_kernel_motor.predict(X_motor_test_MLKR)

    r2_train_motor = r2_score(y_motor_train_kf, y_train_pred_motor)
    r2_test_motor = r2_score(y_motor_value_kf, y_test_pred_motor)

    X_total_train_PCA = pca_transformations.transform(X_train_kf)
    X_total_test_PCA = pca_transformations.transform(X_value_kf)

    mlkr_total = MLKR(n_components=n, init='pca')
    mlkr_total.fit(X_total_train_PCA, y_total_train_kf)
    X_total_train_MLKR = mlkr_total.transform(X_total_train_PCA)
    X_total_test_MLKR = mlkr_total.transform(X_total_test_PCA)

```

```

gauss_kernel_total = KernelRidge(kernel='rbf')
gauss_kernel_total.fit(X_total_train_MLKR, y_total_train_kf)

y_train_pred_total = gauss_kernel_total.predict(X_total_train_MLKR)
y_test_pred_total = gauss_kernel_total.predict(X_total_test_MLKR)

r2_train_total = r2_score(y_total_train_kf, y_train_pred_total)
r2_test_total = r2_score(y_total_value_kf, y_test_pred_total)

return (n, r2_train_motor, r2_test_motor, r2_train_total, r2_test_total)

def process_m(n, kf, X_train_scaled, y_motor_train, y_total_train):
    results = Parallel(n_jobs=-1)(
        delayed(process_fold)(train_i, value_i, n, X_train_scaled,
↪y_motor_train, y_total_train)
        for train_i, value_i in kf.split(X_train_scaled)
    )

    return results

parallel_results = Parallel(n_jobs=-1)(
    delayed(process_m)(n, kf, X_train_scaled, y_motor_train, y_total_train) for
↪n in [n]
)

for result in parallel_results:
    for fold_result in result:
        print(f"Components: {fold_result[0]}, R^2 Train Motor:
↪{fold_result[1]}, R^2 Test Motor: {fold_result[2]}, R^2 Train Total:
↪{fold_result[3]}, R^2 Test Total: {fold_result[4]}")

```

```

[70]: import pandas as pd
from sklearn.model_selection import train_test_split, KFold
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.kernel_ridge import KernelRidge
from sklearn.metrics import r2_score
from metric_learn import MLKR
from joblib import Parallel, delayed

path = './parkinsons_updrs.data'

data = pd.read_csv(path)

X = data.drop(columns=['subject#', 'motor_UPDRS', 'total_UPDRS'])
y_motor = data['motor_UPDRS']

```

```

y_total = data['total_UPDRS']

X_train, X_test, y_motor_train, y_motor_test, y_total_train, y_total_test =
    ↪train_test_split(X, y_motor, y_total, test_size=0.3, random_state=42)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

kf = KFold(n_splits=5, shuffle=True, random_state=42)

n = 10

pca_transformations = PCA(n_components=n).fit(X_train_scaled)

def process_fold(train_i, value_i, n, X_train_scaled, y_motor_train,
    ↪y_total_train):
    X_train_kf, X_value_kf = X_train_scaled[train_i], X_train_scaled[value_i]
    y_motor_train_kf, y_motor_value_kf = y_motor_train.iloc[train_i],
    ↪y_motor_train.iloc[value_i]
    y_total_train_kf, y_total_value_kf = y_total_train.iloc[train_i],
    ↪y_total_train.iloc[value_i]

    X_motor_train_PCA = pca_transformations.transform(X_train_kf)
    X_motor_test_PCA = pca_transformations.transform(X_value_kf)

    mlkr_motor = MLKR(n_components=n, init='pca')
    mlkr_motor.fit(X_motor_train_PCA, y_motor_train_kf)
    X_motor_train_MLKR = mlkr_motor.transform(X_motor_train_PCA)
    X_motor_test_MLKR = mlkr_motor.transform(X_motor_test_PCA)

    gauss_kernel_motor = KernelRidge(kernel='rbf')
    gauss_kernel_motor.fit(X_motor_train_MLKR, y_motor_train_kf)

    y_train_pred_motor = gauss_kernel_motor.predict(X_motor_train_MLKR)
    y_test_pred_motor = gauss_kernel_motor.predict(X_motor_test_MLKR)

    r2_train_motor = r2_score(y_motor_train_kf, y_train_pred_motor)
    r2_test_motor = r2_score(y_motor_value_kf, y_test_pred_motor)

    X_total_train_PCA = pca_transformations.transform(X_train_kf)
    X_total_test_PCA = pca_transformations.transform(X_value_kf)

    mlkr_total = MLKR(n_components=n, init='pca')
    mlkr_total.fit(X_total_train_PCA, y_total_train_kf)
    X_total_train_MLKR = mlkr_total.transform(X_total_train_PCA)
    X_total_test_MLKR = mlkr_total.transform(X_total_test_PCA)

```

```

gauss_kernel_total = KernelRidge(kernel='rbf')
gauss_kernel_total.fit(X_total_train_MLKR, y_total_train_kf)

y_train_pred_total = gauss_kernel_total.predict(X_total_train_MLKR)
y_test_pred_total = gauss_kernel_total.predict(X_total_test_MLKR)

r2_train_total = r2_score(y_total_train_kf, y_train_pred_total)
r2_test_total = r2_score(y_total_value_kf, y_test_pred_total)

return (n, r2_train_motor, r2_test_motor, r2_train_total, r2_test_total)

def process_m(n, kf, X_train_scaled, y_motor_train, y_total_train):
    results = Parallel(n_jobs=-1)(
        delayed(process_fold)(train_i, value_i, n, X_train_scaled,
↪y_motor_train, y_total_train)
        for train_i, value_i in kf.split(X_train_scaled)
    )

    return results

parallel_results = Parallel(n_jobs=-1)(
    delayed(process_m)(n, kf, X_train_scaled, y_motor_train, y_total_train) for
↪n in [n]
)

for result in parallel_results:
    for fold_result in result:
        print(f"Components: {fold_result[0]}, R^2 Train Motor:
↪{fold_result[1]}, R^2 Test Motor: {fold_result[2]}, R^2 Train Total:
↪{fold_result[3]}, R^2 Test Total: {fold_result[4]}")

```

```

-----
KeyboardInterrupt                                Traceback (most recent call last)
Cell In[70], line 83
    76     results = Parallel(n_jobs=-1)(
    77         delayed(process_fold)(train_i, value_i, n, X_train_scaled,
↪y_motor_train, y_total_train)
    78         for train_i, value_i in kf.split(X_train_scaled)
    79     )
    81     return results
---> 83 parallel_results = Parallel(n_jobs=-1)(
    84     delayed(process_m)(n, kf, X_train_scaled, y_motor_train,
↪y_total_train) for n in [n]
    85 )
    87 # Print results for each fold
    88 for result in parallel_results:

```

```
File ~/anaconda3/envs/EE559/lib/python3.11/site-packages/joblib/parallel.py:
-> 2007, in Parallel.__call__(self, iterable)
    2001 # The first item from the output is blank, but it makes the interpreter
    2002 # progress until it enters the Try/Except block of the generator and
    2003 # reaches the first `yield` statement. This starts the asynchronous
    2004 # dispatch of the tasks to the workers.
    2005 next(output)
-> 2007 return output if self.return_generator else list(output)
```

```
File ~/anaconda3/envs/EE559/lib/python3.11/site-packages/joblib/parallel.py:
-> 1650, in Parallel._get_outputs(self, iterator, pre_dispatch)
    1647     yield
    1649     with self._backend.retrieval_context():
-> 1650         yield from self._retrieve()
    1652 except GeneratorExit:
    1653     # The generator has been garbage collected before being fully
    1654     # consumed. This aborts the remaining tasks if possible and warn
    1655     # the user if necessary.
    1656     self._exception = True
```

```
File ~/anaconda3/envs/EE559/lib/python3.11/site-packages/joblib/parallel.py:
-> 1762, in Parallel._retrieve(self)
    1757 # If the next job is not ready for retrieval yet, we just wait for
    1758 # async callbacks to progress.
    1759 if ((len(self._jobs) == 0) or
    1760     (self._jobs[0].get_status(
    1761         timeout=self.timeout) == TASK_PENDING)):
-> 1762     time.sleep(0.01)
    1763     continue
    1765 # We need to be careful: the job list can be filling up as
    1766 # we empty it and Python list are not thread-safe by
    1767 # default hence the use of the lock
```

KeyboardInterrupt:

```
[ ]: import pandas as pd
from sklearn.model_selection import train_test_split, KFold
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.kernel_ridge import KernelRidge
from sklearn.metrics import r2_score
from metric_learn import MLKR
from joblib import Parallel, delayed

path = './parkinsons_updrs.data'
```

```

data = pd.read_csv(path)

X = data.drop(columns=['subject#', 'motor_UPDRS', 'total_UPDRS'])
y_motor = data['motor_UPDRS']
y_total = data['total_UPDRS']

X_train, X_test, y_motor_train, y_motor_test, y_total_train, y_total_test = \
    ↪train_test_split(X, y_motor, y_total, test_size=0.3, random_state=42)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

kf = KFold(n_splits=5, shuffle=True, random_state=42)

n = 15

pca_transformations = PCA(n_components=n).fit(X_train_scaled)

def process_fold(train_i, value_i, n, X_train_scaled, y_motor_train, \
    ↪y_total_train):
    X_train_kf, X_value_kf = X_train_scaled[train_i], X_train_scaled[value_i]
    y_motor_train_kf, y_motor_value_kf = y_motor_train.iloc[train_i], \
    ↪y_motor_train.iloc[value_i]
    y_total_train_kf, y_total_value_kf = y_total_train.iloc[train_i], \
    ↪y_total_train.iloc[value_i]

    X_motor_train_PCA = pca_transformations.transform(X_train_kf)
    X_motor_test_PCA = pca_transformations.transform(X_value_kf)

    mlkr_motor = MLKR(n_components=n, init='pca')
    mlkr_motor.fit(X_motor_train_PCA, y_motor_train_kf)
    X_motor_train_MLKR = mlkr_motor.transform(X_motor_train_PCA)
    X_motor_test_MLKR = mlkr_motor.transform(X_motor_test_PCA)

    gauss_kernel_motor = KernelRidge(kernel='rbf')
    gauss_kernel_motor.fit(X_motor_train_MLKR, y_motor_train_kf)

    y_train_pred_motor = gauss_kernel_motor.predict(X_motor_train_MLKR)
    y_test_pred_motor = gauss_kernel_motor.predict(X_motor_test_MLKR)

    r2_train_motor = r2_score(y_motor_train_kf, y_train_pred_motor)
    r2_test_motor = r2_score(y_motor_value_kf, y_test_pred_motor)

    X_total_train_PCA = pca_transformations.transform(X_train_kf)
    X_total_test_PCA = pca_transformations.transform(X_value_kf)

```

```

mlkr_total = MLKR(n_components=n, init='pca')
mlkr_total.fit(X_total_train_PCA, y_total_train_kf)
X_total_train_MLKR = mlkr_total.transform(X_total_train_PCA)
X_total_test_MLKR = mlkr_total.transform(X_total_test_PCA)

gauss_kernel_total = KernelRidge(kernel='rbf')
gauss_kernel_total.fit(X_total_train_MLKR, y_total_train_kf)

y_train_pred_total = gauss_kernel_total.predict(X_total_train_MLKR)
y_test_pred_total = gauss_kernel_total.predict(X_total_test_MLKR)

r2_train_total = r2_score(y_total_train_kf, y_train_pred_total)
r2_test_total = r2_score(y_total_value_kf, y_test_pred_total)

return (n, r2_train_motor, r2_test_motor, r2_train_total, r2_test_total)

def process_m(n, kf, X_train_scaled, y_motor_train, y_total_train):
    results = Parallel(n_jobs=-1)(
        delayed(process_fold)(train_i, value_i, n, X_train_scaled,
↪y_motor_train, y_total_train)
        for train_i, value_i in kf.split(X_train_scaled)
    )

    return results

parallel_results = Parallel(n_jobs=-1)(
    delayed(process_m)(n, kf, X_train_scaled, y_motor_train, y_total_train) for
↪n in [n]
)

for result in parallel_results:
    for fold_result in result:
        print(f"Components: {fold_result[0]}, R^2 Train Motor:
↪{fold_result[1]}, R^2 Test Motor: {fold_result[2]}, R^2 Train Total:
↪{fold_result[3]}, R^2 Test Total: {fold_result[4]}")

```

```

[ ]: import pandas as pd
from sklearn.model_selection import train_test_split, KFold
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.kernel_ridge import KernelRidge
from sklearn.metrics import r2_score
from metric_learn import MLKR
from joblib import Parallel, delayed

path = './parkinsons_updrs.data'

```



```

data = pd.read_csv(path)

X = data.drop(columns=['subject#', 'motor_UPDRS', 'total_UPDRS'])
y_motor = data['motor_UPDRS']
y_total = data['total_UPDRS']

X_train, X_test, y_motor_train, y_motor_test, y_total_train, y_total_test = ␣
    ↪ train_test_split(X, y_motor, y_total, test_size=0.3, random_state=42)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

kf = KFold(n_splits=5, shuffle=True, random_state=42)

n = 19

pca_transformations = PCA(n_components=n).fit(X_train_scaled)

def process_fold(train_i, value_i, n, X_train_scaled, y_motor_train, ␣
    ↪ y_total_train):
    X_train_kf, X_value_kf = X_train_scaled[train_i], X_train_scaled[value_i]
    y_motor_train_kf, y_motor_value_kf = y_motor_train.iloc[train_i], ␣
    ↪ y_motor_train.iloc[value_i]
    y_total_train_kf, y_total_value_kf = y_total_train.iloc[train_i], ␣
    ↪ y_total_train.iloc[value_i]

    mlkr_motor = MLKR(n_components=n, init='auto')
    mlkr_motor.fit(X_train_kf, y_motor_train_kf)
    X_motor_train_MLKR = mlkr_motor.transform(X_train_kf)
    X_motor_test_MLKR = mlkr_motor.transform(X_value_kf)

    gauss_kernel_motor = KernelRidge(kernel='rbf')
    gauss_kernel_motor.fit(X_motor_train_MLKR, y_motor_train_kf)

    y_train_pred_motor = gauss_kernel_motor.predict(X_motor_train_MLKR)
    y_test_pred_motor = gauss_kernel_motor.predict(X_motor_test_MLKR)

    r2_train_motor = r2_score(y_motor_train_kf, y_train_pred_motor)
    r2_test_motor = r2_score(y_motor_value_kf, y_test_pred_motor)

    mlkr_total = MLKR(n_components=n, init='auto')
    mlkr_total.fit(X_train_kf, y_total_train_kf)
    X_total_train_MLKR = mlkr_total.transform(X_train_kf)
    X_total_test_MLKR = mlkr_total.transform(X_value_kf)

```

```

gauss_kernel_total = KernelRidge(kernel='rbf')
gauss_kernel_total.fit(X_total_train_MLKR, y_total_train_kf)

y_train_pred_total = gauss_kernel_total.predict(X_total_train_MLKR)
y_test_pred_total = gauss_kernel_total.predict(X_total_test_MLKR)

r2_train_total = r2_score(y_total_train_kf, y_train_pred_total)
r2_test_total = r2_score(y_total_value_kf, y_test_pred_total)

return (n, r2_train_motor, r2_test_motor, r2_train_total, r2_test_total)

def process_m(n, kf, X_train_scaled, y_motor_train, y_total_train):
    results = Parallel(n_jobs=-1)(
        delayed(process_fold)(train_i, value_i, n, X_train_scaled,
↪y_motor_train, y_total_train)
        for train_i, value_i in kf.split(X_train_scaled)
    )

    return results

parallel_results = Parallel(n_jobs=-1)(
    delayed(process_m)(n, kf, X_train_scaled, y_motor_train, y_total_train) for
↪n in [n]
)

for result in parallel_results:
    for fold_result in result:
        print(f"Components: {fold_result[0]}, R^2 Train Motor:
↪{fold_result[1]}, R^2 Test Motor: {fold_result[2]}, R^2 Train Total:
↪{fold_result[3]}, R^2 Test Total: {fold_result[4]}")

```

- (c) Use sklearn's neural network implementation to train a neural network with two outputs that predicts motor UPDRS and total UPDRS. Use a single layer. You are responsible to determine other architectural parameters of the network, including the number of neurons in the hidden and output layers, method of optimization, type of activation functions, and the L2 "regularization" parameter etc. You should determine the design parameters via trial and error, by testing your trained network on the test set and choosing the architecture that yields the smallest test error. For this part, set early-stopping=False. Remember to standardize your features. Report your R^2 on both training and test sets. (20 pts)

```

[91]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import r2_score
from sklearn.neural_network import MLPRegressor

```

```

path = './parkinsons_updrs.data'
data = pd.read_csv(path)

X = data.drop(columns=['subject#', 'motor_UPDRS', 'total_UPDRS'])
y = data[['motor_UPDRS', 'total_UPDRS']]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
    ↪random_state=42)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

nn = MLPRegressor(hidden_layer_sizes=(200,), activation='tanh', solver='sgd',
    ↪alpha=0.001, max_iter=5000, learning_rate_init=0.001,
    ↪learning_rate='adaptive', random_state=42, early_stopping=False,
    ↪validation_fraction=0.1)

nn.fit(X_train_scaled, y_train)

y_train_pred = nn.predict(X_train_scaled)
r2_train_motor = r2_score(y_train['motor_UPDRS'], y_train_pred[:, 0])
r2_train_total = r2_score(y_train['total_UPDRS'], y_train_pred[:, 1])

y_test_pred = nn.predict(X_test_scaled)
r2_test_motor = r2_score(y_test['motor_UPDRS'], y_test_pred[:, 0])
r2_test_total = r2_score(y_test['total_UPDRS'], y_test_pred[:, 1])

print(f'R^2 on training set for motor UPDRS: {r2_train_motor}')
print(f'R^2 on training set for total UPDRS: {r2_train_total}')
print(f'R^2 on test set for motor UPDRS: {r2_test_motor}')
print(f'R^2 on test set for total UPDRS: {r2_test_total}')

```

R² on training set for motor UPDRS: 0.9591315893397608

R² on training set for total UPDRS: 0.9641449844842482

R² on test set for motor UPDRS: 0.8292075924024247

R² on test set for total UPDRS: 0.8348681095343674

(d) Use the design parameters that you chose in the first part and train a neural network, but this time set early-stopping=True. Research what early stopping is, and compare the performance of your network on the test set with the previous network. You can leave the validation-fraction as the default (0.1) or change it to see whether you can obtain a better model. Remember to standardize your features. Report your R² on both training and test sets. (10 pts)

```

[86]: import numpy as np
import pandas as pd

```

```

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import r2_score
from sklearn.neural_network import MLPRegressor

path = './parkinsons_updrs.data'
data = pd.read_csv(path)

X = data.drop(columns=['subject#', 'motor_UPDRS', 'total_UPDRS'])
y = data[['motor_UPDRS', 'total_UPDRS']]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
    ↪random_state=42)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

nn = MLPRegressor(hidden_layer_sizes=(200,), activation='tanh', solver='sgd',
    ↪alpha=0.0001, max_iter=5000, learning_rate_init=0.01,
    ↪learning_rate='adaptive', random_state=42, early_stopping=True,
    ↪validation_fraction=0.1)

nn.fit(X_train_scaled, y_train)

y_train_pred = nn.predict(X_train_scaled)
r2_train_motor = r2_score(y_train['motor_UPDRS'], y_train_pred[:, 0])
r2_train_total = r2_score(y_train['total_UPDRS'], y_train_pred[:, 1])

y_test_pred = nn.predict(X_test_scaled)
r2_test_motor = r2_score(y_test['motor_UPDRS'], y_test_pred[:, 0])
r2_test_total = r2_score(y_test['total_UPDRS'], y_test_pred[:, 1])

print(f'R^2 on training set for motor UPDRS: {r2_train_motor}')
print(f'R^2 on training set for total UPDRS: {r2_train_total}')
print(f'R^2 on test set for motor UPDRS: {r2_test_motor}')
print(f'R^2 on test set for total UPDRS: {r2_test_total}')

```

R² on training set for motor UPDRS: 0.9321351544702516

R² on training set for total UPDRS: 0.9338009724721683

R² on test set for motor UPDRS: 0.8119978743947998

R² on test set for total UPDRS: 0.8058333276066811

With tuning I was able to get R^2 values in the 0.80s range, but still worse than without early stopping